



Iskanje brezna velikih podatkov v realnem času

Big Data Analytics in Real-time

© 2025 Simon Šanca, simon.sanca@uib.no

[University of Bergen / IT-Platform, Team Linux](#)

OpenSearch specifikacija

- **Predstavitev:** O'Reilly Emerging Technology Conference - 15. marca 2005.
- **Prvotni cilj:**
 - Standardizacija objave rezultatov iskanja z uporabo opisnih datotek v XML,
 - standardizacija poizvedb ter odzivov v RSS ali Atom (query syntax).
- **Uporabnost:**
 - Hitro iskanje spletnih strani in enostavno deljenje rezultatov.
 - *A way for websites and search engines to publish search results in a standard and accessible format.*

OpenSearch je postal široko podprt v brskalnikih, kot so Firefox, Safari in Chrome, kar je takrat omogočilo dodajanje prilagojenih iskalnikov v iskalno vrstico brskalnika.

OpenSearch Project

[OpenSearch](#) is a community-driven, open-source search and analytics suite used by developers to ingest, search, visualize, and analyze data.

- veja Elasticsearch in Kibane, ki sta plačljiva.

v1.0 - julij 2021; Apache Licence, Version 2.0

Glavne komponente:

- **OpenSearch:** Shramba podatkov in iskalnik za hitro obdelavo poizvedb.
- **OpenSearch Dashboards:** Orodje za vizualizacijo podatkov in up. vmesnik.
- **OpenSearch Data Prepper:** Strežniški zbiralnik podatkov za pripravo podatkov.

Uporabniki lahko razširijo funkcionalnost OpenSearch z izbiro vtičnikov, ki izboljšajo iskanje, analitiko, opazovanje, varnost, strojno učenje in še več.

OpenSearch danes

Dokumentacija: <https://opensearch.org/docs/latest/>

Repozitorij: <https://github.com/opensearch-project>

Do avgusta 2024 je AWS poročal o "desetih tisočih" strankah, z več kot 700 milijoni prenosov in prispevki tisočev razvijalcev.

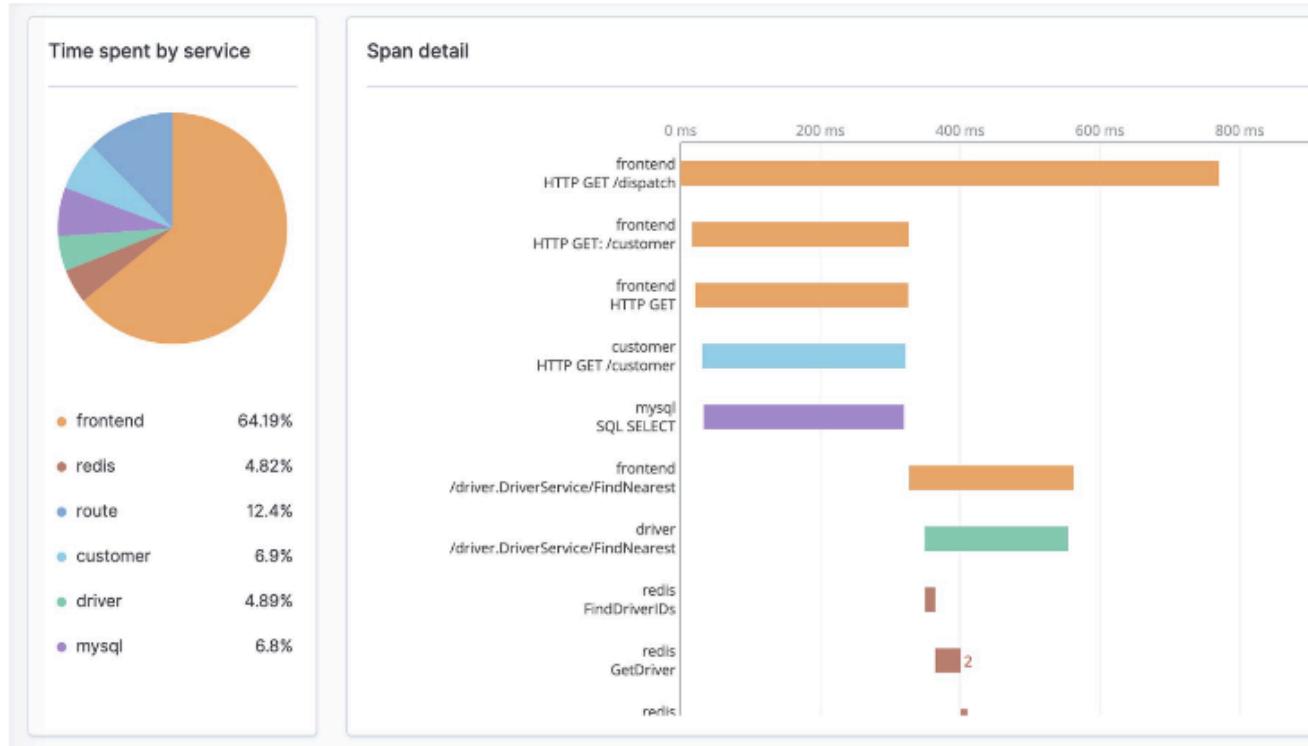
Septembra 2024 se je lastništvo preneslo z AWS na **OpenSearch Software Foundation** pod okriljem **Linux Foundation**.

Gre za porazdeljen iskalni in analitič sistem, ki temelji na vektorski podatkovni bazi.

Application Performance Monitoring, Log Analytics, Big Data Analytics, Time Series Analysis, Data Visualization in [ostalo](#).



Projekt DataOPS: Log Analytics; normalizacija / standardizacija zapisov (logs)



Logs

```

9  61.159.121.13 - - [10/Apr/2018:14:08:06 +0200] "GET /tipps/google-suche.html HTTP/1.1" 200 6731 "https://www.sti
10 61.159.121.13 - - [10/Apr/2018:14:08:06 +0200] "GET /styles/format.css HTTP/1.1" 200 6831 "http://www.sti
11 61.159.121.13 - [10/Apr/2018:14:08:06 +0200] "GET /stylies/formatxp.css HTTP/1.1" 200 XSL "http://www.sti
12 61.159.121.13 - - [10/Apr/2018:14:08:06 +0200] "GET /tipps/such-vorschlaege.png HTTP/1.1" 200 8164 "http:
13 61.159.121.13 - - [10/Apr/2018:14:08:06 +0200] "GET /stylies/kontakt.png HTTP/1.1" 200 898 "http://www.sti
14 61.159.121.13 - - [10/Apr/2018:14:08:06 +0200] "GET /tipps/muchbegrieff-operator.png HTTP/1.1" 200 X660 "h
15 61.159.121.13 - - [10/Apr/2018:14:08:06 +0200] "GET /tipps/google-autosuggest-styling.png HTTP/1.1" 20
16 61.159.121.13 - - [10/Apr/2018:14:08:06 +0200] "GET /emglisch-30.png HTTP/1.1" 200 1035 "http://www.sti
17 61.159.121.13 - - [10/Apr/2018:14:08:06 +0200] "GET /imgo.png HTTP/1.1" 200 6006 "http://www.stichpunkt.d
18 61.159.121.13 - - [10/Apr/2018:14:08:06 +0200] "GET /favicon.ico HTTP/1.1" 200 3262 "-" "Mozilla/5.0 (Win
19 61.159.121.13 - - [10/Apr/2018:14:08:06 +0200] "GET /favicon.ico HTTP/1.1" 499 0 "-" "Mozilla/5.0 (Windows
20 2003:6fb121:6e57:3034:770e:2d1:131d - - [10/Apr/2018:14:08:06 +0200] "GET /styles/format.css HTTP/1.1"
21 2003:6fb121:6e57:3034:770e:2d1:131d - - [10/Apr/2018:14:08:26 +0200] "GET /stylies/background-header.png
22 2003:6fb121:6e57:3034:770e:2d1:131d - - [10/Apr/2018:14:08:26 +0200] "GET /stylies/fonts/openSans-Regular
23 2003:6fb121:6e57:3034:770e:2d1:131d - - [10/Apr/2018:14:08:27 +0200] "GET /abbreviations/internet-acrony
24 2003:6fb121:6e57:3034:770e:2d1:131d - - [10/Apr/2018:14:08:27 +0200] "GET /abbreviations/text-messages.pn
25 2003:6fb121:6e57:3034:770e:2d1:131d - - [10/Apr/2018:14:08:27 +0200] "GET /stylies/contact.css HTTP/1.1"
26 2003:6fb121:6e57:3034:770e:2d1:131d - - [10/Apr/2018:14:08:27 +0200] "GET /stylies/format.css HTTP/1.1"
27 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-content/themes/hueman/style.css?ver=4.9.3 HU
28 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-content/themes/hueman/style.css?ver=4.9.3 HU
29 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-content/themes/hueman/assets/front/css/main.css
30 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-content/themes/hueman-child/style.css?ver=3.3.25
31 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-content/themes/hueman/assets/front/css/font-aweso
32 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-includes/js/jquery/jquery.js?ver=1.12.4 HTTP/1.1"
33 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-includes/js/jquery/jquery-migrate.min.js?ver=1.4.
34 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-content/themes/hueman-child/custom.js?ver=1 HTTP/
35 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-content/plugins/hueman-addons/assets/front
36 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-includes/js/underscore.min.js?ver=1.8.0 HTTP/1.1"
37 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-includes/js/wp-tables.min.js?ver=4.9.3 HTTP/1.1"
38 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-content/themes/hueman/assets/front/js/scripts.min
39 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-includes/js/wp-emoji-release.min.js?ver=4.9.3 HTTP
40 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-content/themes/hueman/assets/front/img/slidesba
41 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /wp-content/themes/hueman/assets/front/img/slidesba
42 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /zischen/rvuzusla-minzalne-nots.mp3 HTTP/1.1" 304 0 "
43 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /zischen/rvuzusla-hazmonische-toono-1-bis-3.mp3 HTTP/
44 63.172.164.99 - - [10/Apr/2018:17:44:54 +0200] "GET /zischen/rvuzusla-fussball-stadion.mp3 HTTP/1.1" 304

```

OpenSearch Dashboards

- Da vemo kaj se dogaja z našimi serverji.



Kako hranimo podatke v OpenSearch?

Dokument / Document

- Dokument je enota, ki shranjuje informacije (besedilo ali strukturirane podatke). V OpenSearch so dokumenti shranjeni v formatu JSON.
- En dokument pomeni en zapis (vrstico) v podatkovni bazi.
- Ko iščemo informacije, OpenSearch vrne dokumente, povezane z našim iskanjem.

```
{  
  "ime": "Jurka",  
  "ocena": 5.0,  
  "leto_pridelave": 2020  
  "kraj_pridelave": "Središče",  
  "smo_spili": true  
}
```

Indeks / Index

- Zbirka dokumentov / a collection of documents.
- V relacijski podatkovni bazi bi indeks predstavljal tabelo.
- Ko iščemo, vbistvu poizvedujemo po podatkih, ki jih hrani indeks.

In vino veritas:

Ime	Ocena	Leto pridelave	Kraj pridelave	Smo spili
Jurka	5.0	2020	Središče	Da
Refošk	4.6	2019	Koper	Ne
Modra frankinja	4.9	2021	Bizejjsko	Da

Primer indeksa v OpenSearch

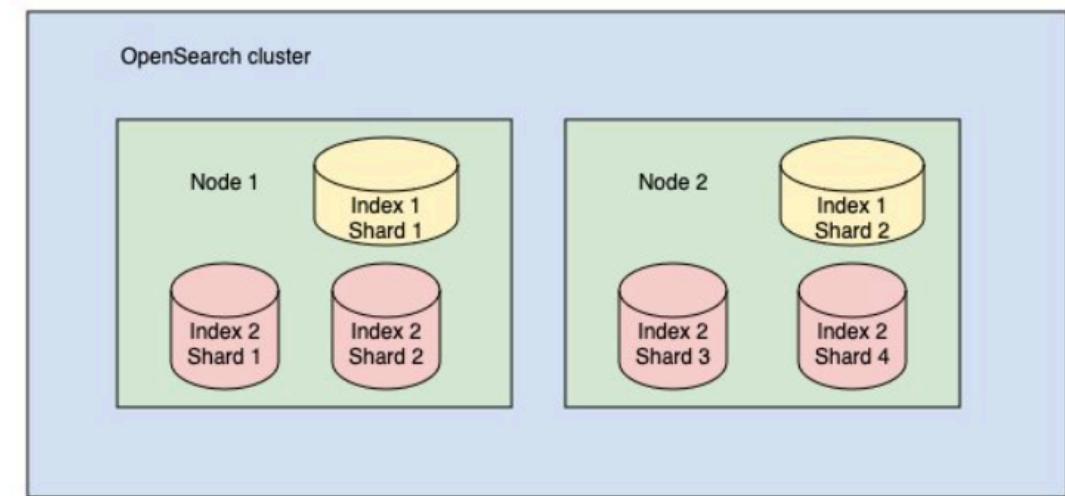
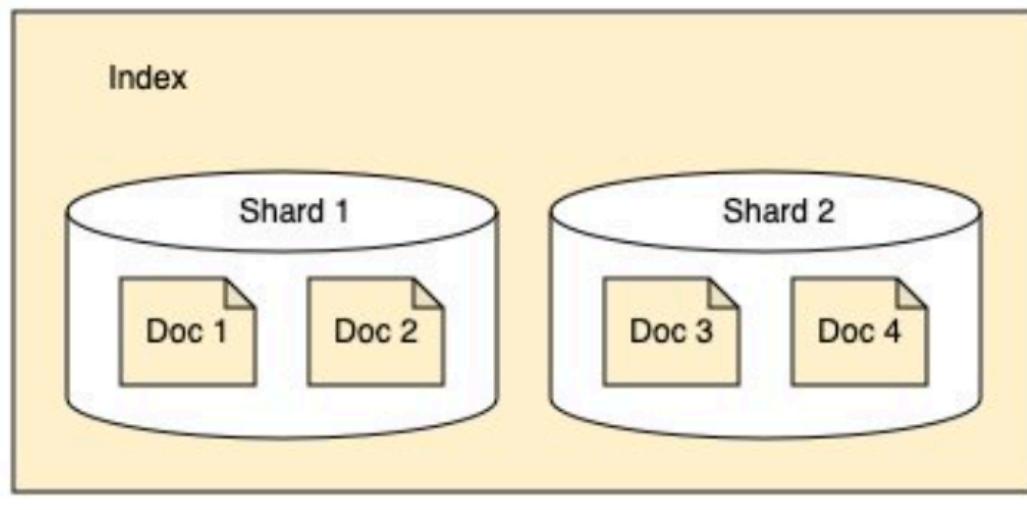
```
{  
  "index": {  
    "_index": "vina"  
  },  
  "data": [  
    {  
      "ime": "Jurka",  
      "ocena": 5.0,  
      "leto_pridelave": 2020,  
      "kraj_pridelave": "Središče",  
      "smo_spili": true  
    },  
    {  
      "ime": "Refošk",  
      "ocena": 4.6,  
      "leto_pridelave": 2019,  
      "kraj_pridelave": "Koper",  
      "smo_spili": false  
    },  
    {  
      "ime": "Modra frankinja",  
      "ocena": 4.9,  
      "leto_pridelave": 2021,  
      "kraj_pridelave": "Bizejjsko",  
      "smo_spili": true  
    }  
  ]  
}
```

Gruče in vozli / Clusters and nodes

- OpenSearch je zasnovan kot porazdeljeni iskalnik; lahko deluje na enem ali več strežnikih (vozliščih), ki shranjujejo podatke in obdelujejo poizvedbe.
- Gruča (cluster) - je zbirka vozlišč / vozlov.
- V gruči z enim samim vozliščem mora ena naprava opraviti vse naloge: upravljati stanje gruče, indeksirati, predobdelati podatke pred indeksiranjem, ipd.
- Ker s podatki gruča raste, lahko naloge razdelimo na več vozlov in poskrbimo za nemoteno delovanje sistema na večih strežnikih.
- **Cluster manager node** - usklajuje operacije na ravni gruče, kot npr. ustvarjanje indeksa. Vozli med seboj komunicirajo; več vozlišč, hitre poizvedbe.

Razseki / Shards

OpenSearch razdeli indekse na razseke. Vsak razsek shranjuje podmnožico vseh dokumentov znotraj indeksa.

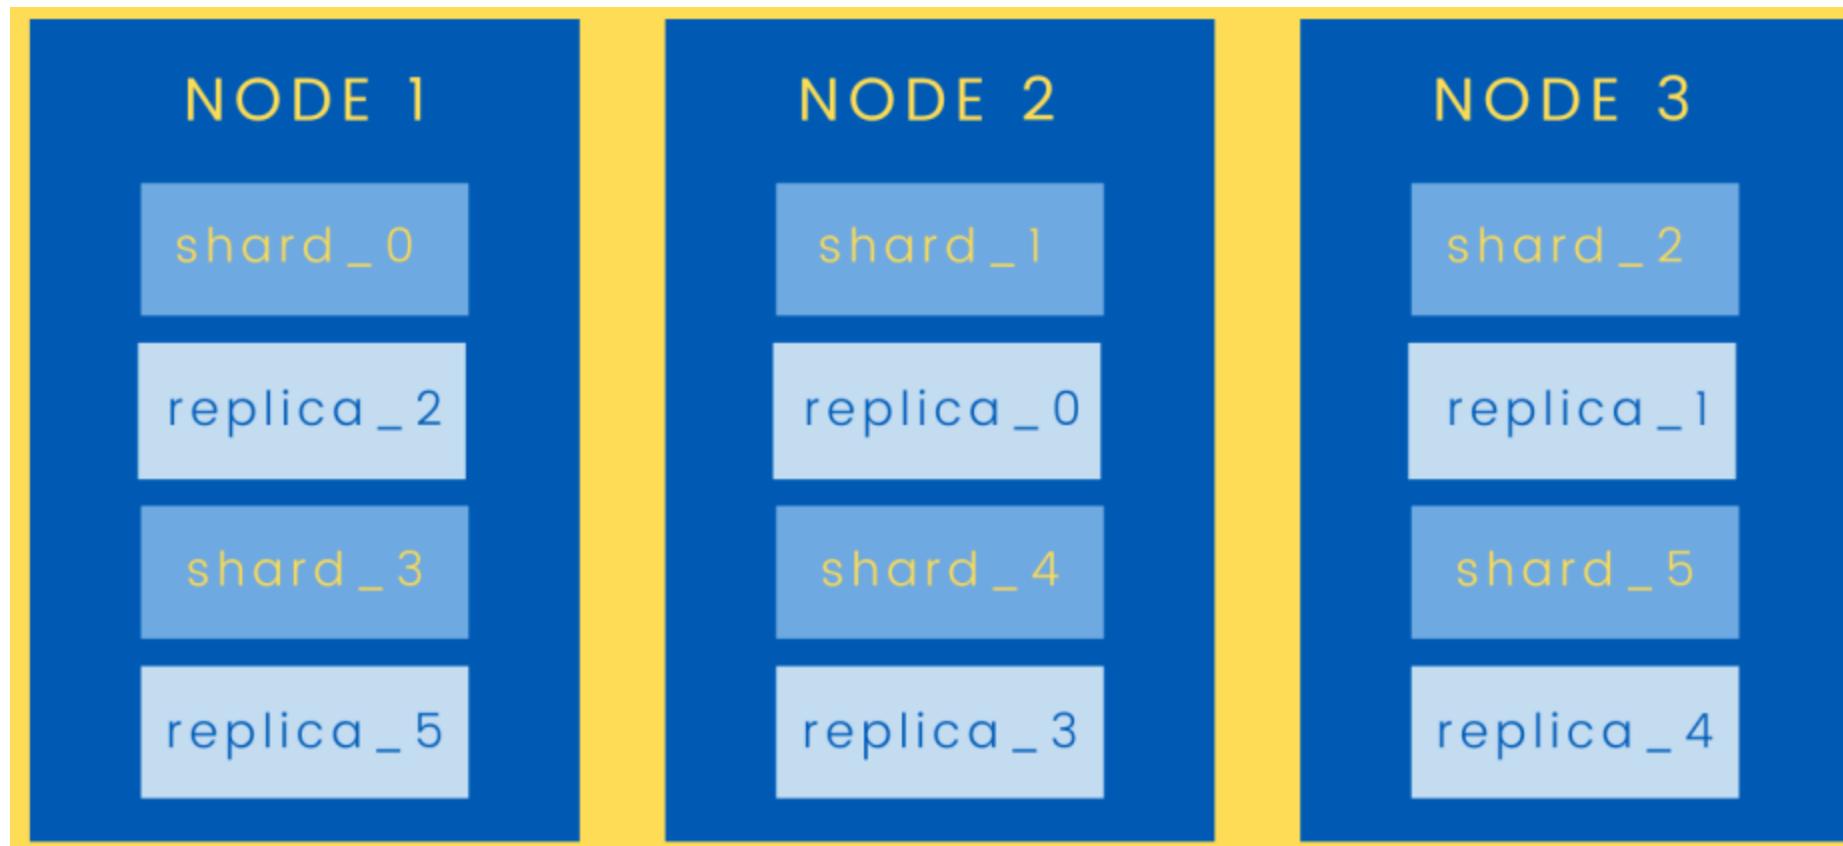


Razseki se uporabljajo za enakomerno porazdeljevanje po vozliščih v gruči.

Primer: Indeks velikosti 400 GB je morda prevelik, da bi ga eno samo vozlišče v clustru obvladalo, ampak, če ga razdelimo na 10 razsekov po 40 GB je vse lažje.

Replike / Replicas

Razsek je primarni (primary) ali replika (replica). OpenS. ustvari replika razsek za vsak primarni razsek. Če indeks razdelimo na 10 razsekov, OpenSearch ustvari 10 replik.



Indeksne predloge / Index templates

- Omogočajo preslikavo novih indeksov z vnaprej določenimi nastavitevami.

```
PUT _index_template/vina
{
  "index_patterns": ["vina*"],
  "template": {
    "settings": {
      "number_of_shards": 1,
      "number_of_replicas": 1
    },
    "mappings": {
      "properties": {
        "ime": { "type": "text" },
        "ocena": { "type": "float" },
        "leto_pridelave": { "type": "integer" },
        "kraj_pridelave": { "type": "keyword" },
        "smo_spili": { "type": "boolean" }
      }
    }
  }
}
```

OpenSearch in GEO

Geopoint

- točka s koordinatami: (lat, lon): { "point": { "lat": 40.71, "lon": 74.00 } }

Geoshape

- OS razdeli poligon v trikotniško mrežo in shrani vsak trikotnik v BKD drevo.
- Formati: GeoJSON in WKT.
- Podatkovni tipi:
 - point, multipoint,
 - linestring, multilinestring,
 - polygon, multipolygon
 - geometrycollection, envelope

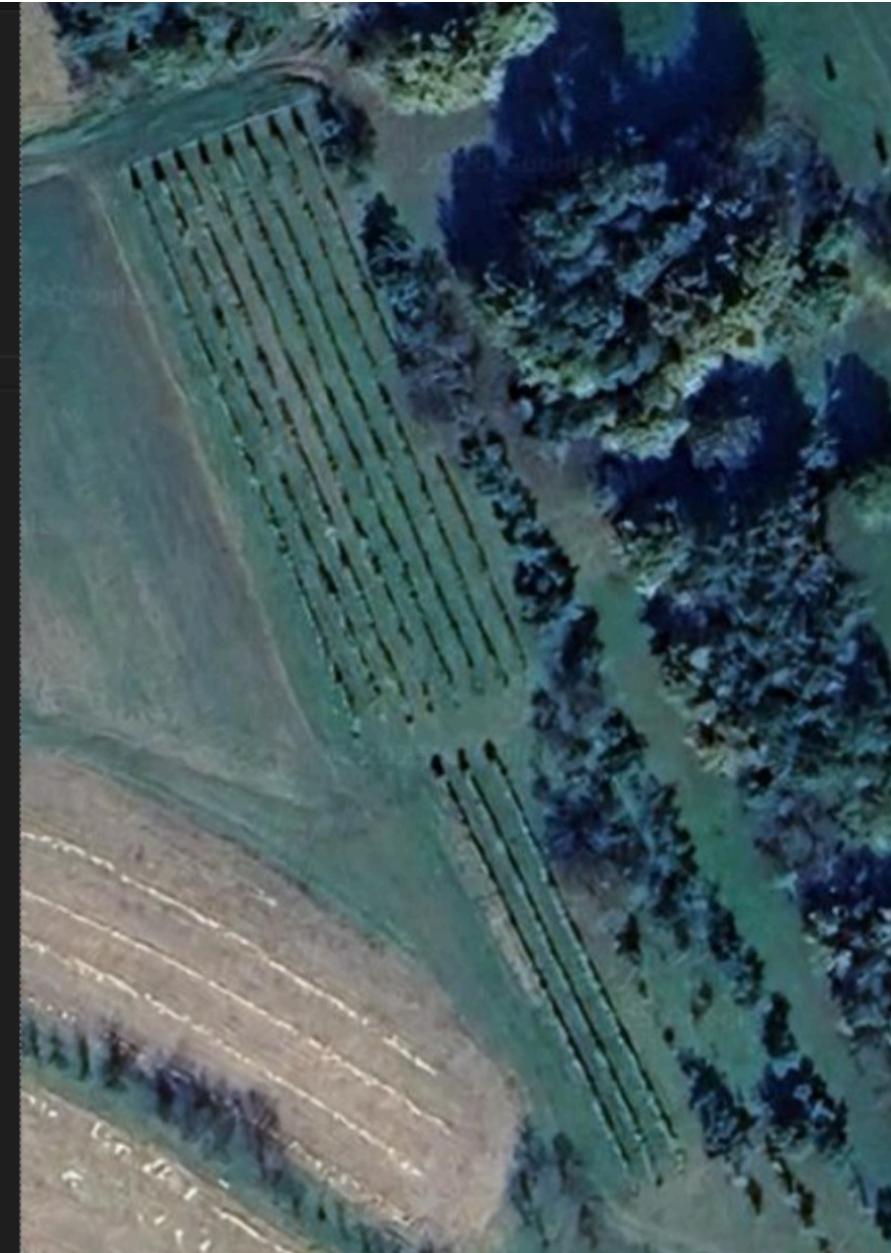
Cartesian field types

- Indeksirajo in iščejo 2D točke na ravnini, ne na sferi. So hitrejše za razvrščanje po razdalji kot geografski tipi. Natančnost: float z eno decimalko.
- Dva tipa:
 - [xy_point](#)
 - [xy_shape](#)

```
{ "point": { "x": 0.5, "y": 4.5 } }
```

```
{"location": {"type": "linestring", "coordinates": [[0.5, 4.5], [-1.5, 2.3]]}}
```

```
{ vinograd.json > {} _source > {} location
1  {
2      "_index": "slovenia_vineyards",
3      "_id": "vinogradi_87_sredisce_06",
4      "_version": 1,
5      "_source": {
6          "name": "Vinogradi",
7          "type": "Vineyard",
8          "location": {
9              "lat": 46.7716,
10             "lon": 16.3060
11         },
12         "area": {
13             "type": "polygon",
14             "coordinates": [
15                 [
16                     [16.305551, 46.772072],
17                     [16.305920, 46.771487],
18                     [16.305974, 46.771508],
19                     [16.306275, 46.771103],
20                     [16.306505, 46.771201],
21                     [16.305829, 46.772140]
22                 ]
23             ],
24             "region": "Goričko",
25             "grape_varieties": ["Jurka", "Isabela", "Modra Frankinja"],
26             "tags": ["vineyard", "wine", "organic", "Središče"],
27             "elevation_m": 250,
28             "annual_production_liters": 1000,
29             "last_updated": "2025-04-08T07:00:00Z"
30         }
31     }
32 }
```



Namestitev / zagon

- Navodila na Githubu.
- <https://opensearch.org/docs/latest/install-and-configure/install-opensearch/docker/>
- YAML in Docker Compose
- `docker compose up -d`
- OOM exception, set: `sysctl -w vm.max_map_count=262144`
- OpenSearch dostopen na: <http://localhost:5601/>
- Uporabniško ime in geslo se lahko dodata v *docker-compose.yml* konfiguracijo.

Our sample file creates two OpenSearch nodes and one OpenSearch Dashboards node with the Security plugin disabled.

```
[root@logstash ~]# vi docker-compose.yml
[root@logstash ~]# docker compose up -d
[+] Running 3/3
  ✓ Container opensearch-node1      Started
    21.2s
  ✓ Container opensearch-node2      Started
    21.2s
  ✓ Container opensearch-dashboards Started
    22.1s
[root@logstash ~]# docker logs opensearch-dashboards
Disabling OpenSearch Security Dashboards Plugin
Removing securityDashboards...
Plugin removal complete
[root@logstash ~]# docker logs opensearch-node1
```

Malo je treba potrpeti in pregledati log datoteke, in počakati da se vse naloži.

```
docker logs opensearch-dashboards
docker logs opensearch-node1
docker logs opensearch-node2
```

<http://localhost:5601/>

The screenshot shows the landing page of the OpenSearch Dashboards interface. At the top center is a white circular logo containing a blue stylized 'S' or gear icon. Below the logo, the text "Welcome to OpenSearch Dashboards" is displayed in a large, dark font. Underneath this text is a decorative graphic featuring various data visualization elements like a bar chart, a pie chart, a line graph, and a table, all set against a background of a small house and a barcode.

Start by adding your data

Add data to your cluster from any source, then analyze and visualize it in real time. Use our solutions to add search anywhere, observe your ecosystem, and protect against security threats.

[Add data](#) [Explore on my own](#)

Preprost Geo primer

Primer geodeta ki z GNSS-jem meri na Goričkem pomeri prikazan vinograd, ker je alternativec ne uporablja Leice, ampak ruski Emlid. Meritve pošilja v pisarno v Mursko Soboto, kjer teče OpenSearch in ustvari preprost Dashboard.

Dober sosed Julius radovedno opazuje geodeta in ga po meritvah povabi pod senco murve na kosilo in na kozarec domače Jurke.

Zgodba bi se lahko začela tudi drugače:

- *Dolnji Slaveči, idilična vas na še bolj idiličnem Goričkem, kjer živijo daleč od tega ponorelega sveta prijazni, predsvem pa marljivi ljudje, ki se znajo skregati ...*

Workflow

1. Uvozimo podatke s curl ali ročno z DevTools v OpenSearch.
2. Preverimo če so podatki naloženi.

```
GET /meritve_sredisce/_search { "query": { "match_all": {} } }
```

3. Iskanje znotaj DevTools je možno z QueryDSL, npr.

```
GET /meritve_sredisce/_search
{
  "query": {
    "term": {
      "fix_quality": "2D"
    }
  }
}
```

4. Definiramo preslikavo podatkovih tipov z **Index Patterns**.
5. Usvarimo si svoj **Dashboard**.
6. **Discover** nam omogoča **iskanje po podatkih**.
7. **Anomaly detection** nam lahko najde anomalije; GNSS meritve s šumom.

Dodatne funkcionalnosti

- Vector search.
- Naučimo si modele strojnega učenja na lastnih podatkih.
- GenAI (RAG, LLMs).

Vprašanja?

Why you so quiet? What's on your mind?



Hvala ekipi OsGeo Slovenija <3



Keep it free and open

OpenSearchCon Europe 2025



- Kdaj? 30. april – 1. maj, Amsterdam
- <https://events.linuxfoundation.org/opensearchcon-europe/>